

Lab 2 - Image Convolution on FPGA, CPUs, and GPUs

ECE 1756 (Fall 2025)

Kuan-Yu Chang (1007359760)

November 04, 2025

1.0 Introduction

This document serves as the lab report for ECE1756 Fall 2025 Section. It is authored by Kuan-Yu Chang (1007359760).

Section 2.0 describes the FPGA implementation of the convolution function, its operation and performance while Section 3.0 describes the same metrics of the CPU and GPU implementation. A comparison of the different implementations will be done in Section 4.0.

2.0 FPGA Implementation

To enable a comparison between FPGA, CPU and GPU, a convolution engine was implemented on an Intel Arria 10 GX device (10AX115N2F45I1SG). The following section describes the details of this implementation as well as performance metrics.

2.1 Convolution and Input Data Details

The FPGA convolution engine processes a 512-pixel-wide image of arbitrary height using a provided 3×3 filter. Pixels arrive in row-major order (top-left first, one pixel per cycle), and a convolution is computed for every 3×3 neighborhood. Because neighborhoods span adjacent rows, incoming pixels must be retained across lines. To do this, we use a serial-in/parallel-out (SIPO) shift register. With a fixed line length of 514 pixels (512 image pixels plus zero padding), the SIPO must be at least 1031 entries long. Pixels shift in from the left and advance one register per clock cycle.

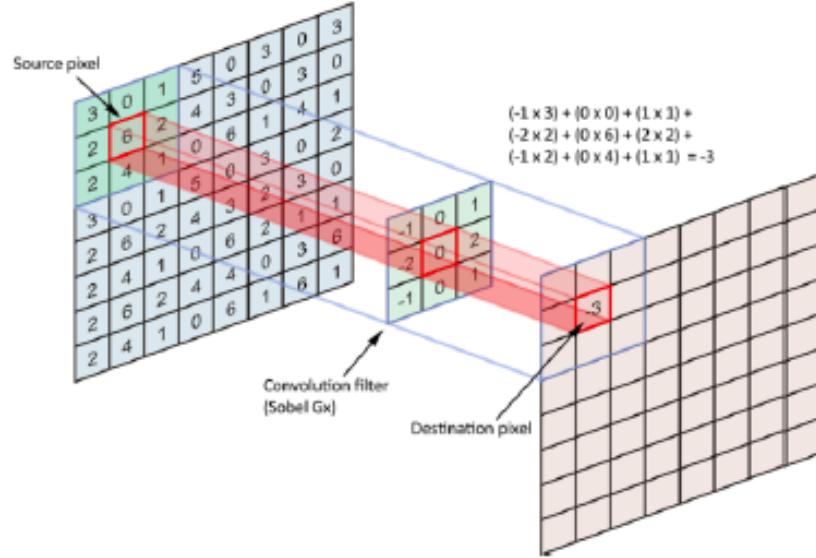


Figure 1: Convolution example of the design

2.2 FPGA Datapath

The pipelined circuit improves performance by dividing the computation into multiple stages, each separated by registers. Each stage performs a smaller portion of the overall calculation, allowing several input values to be processed simultaneously in different stages.

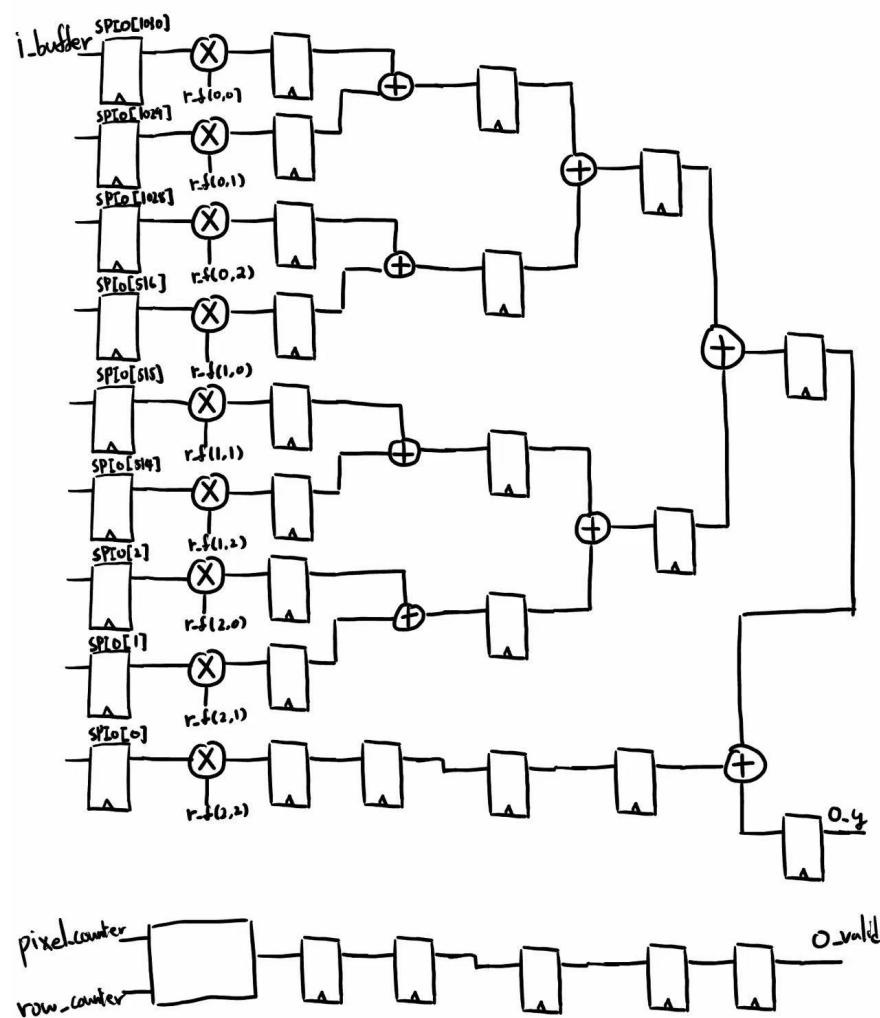


Figure 2: Pipelined convolution datapath

The datapath for this computation of convolution is presented above. Which is composed by 9 multipliers performing the 9 multiplications in parallel than summed in a 3 level adder tree. Data are pipelined at every stage to ensure the data correctness of the output.

2.3 Operation

All items(pixels or paddings) within the convolution window will be a fixed index position in the buffer, and they will be computed with the corresponding filter coefficients. For example, $i_buffer[0]$ will be associated with filter[2][2], $i_buffer[515]$ will be associated with filter[1][1],

and so on. The input data will first be multiplied with the filter window, than added together to preform a pixel. Different filters would give different result to teh convolving an image. As for the o valid signal, pixel counter will determine how many input pixels have already flowed into the buffer(maxium of 1031), and row counter will deal with the offset during row transition. 5 additional pipeline registers are added to match the number of registers in o y path to secure the correctness.

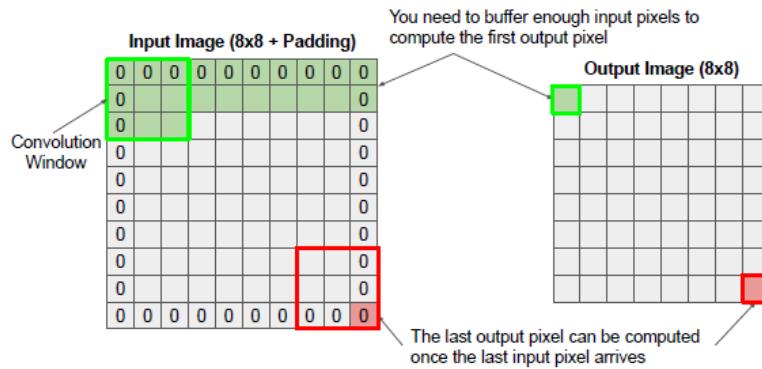


Figure 3: Convolution operation example

2.4 Correctness

The FPGA implementation successfully completes all test cases (Appendix A).

Test	Completed?
1	Yes
2	Yes
3a	Yes
3b	Yes
4a	Yes
4b	Yes
5a	Yes
5b	Yes
6a	Yes
6b	Yes
7a	Yes
7b	Yes

Table 1: FPGA implementation correctness

2.5 Performance Metrics

	Result (Appendix B)
ALM Utilization	2220
DSP Utilization	9
BRAM (M20K) Utilization	0
Maximum Operating Frequency (MHz)	442.28
Cycles for Test 7a (Hinton)	264243
Dynamic Power for one module @ maximum frequency (W)	108.09×10^{-3}
Throughput of one module (GOPS)	7.466
Throughput of a full device (GOPS)	1.251×10^3
Total Power for a full device (W)	12.66

Table 2: FPGA implementation results

Details calculations are included in Appendix B.

3.0 CPU and GPU Implementations

3.1 Operations

Both the CPU and GPU versions compute the image–filter convolution, but they differ in how the work is carried out on their respective hardware. On the CPU, without dedicated convolution hardware, the code relies on nested loops over filters, rows, and columns and, if vectorized, over the filter window to perform multiply-accumulate operations. The GPU version emphasizes memory hierarchy and parallel execution—using CUDA cores, thread blocks, and shared memory—to reduce explicit nesting and improve throughput through better data reuse and concurrency.

3.2 Preformance Metrics

From the table above, we can see that for all the test case given, runtime increases with the number of filters linearly. At every optimization level, the vectorized implementation finishes faster than the basic one—except with O3. The speedup comes from doing fewer explicit nested loops and more efficient arithmetic. With higher compiler optimizations (O2, O3), the compiler applies more aggressive transforms—e.g., deeper loop unrolling and late-stage instruction optimizations—reducing runtime further. The basic version, having more loops, tends to benefit disproportionately from O3, which narrows the gap with the vectorized code. Enabling multithreading yields substantial gains when the number of filters is large, cutting

No. of filters	Runtime (ms) (Appendix C)			
	1	4	16	64
GPU	0.0146194	0.0382884	0.129539	0.494356
CPU (basic - no opt - 1 thread)	6.06443	24.4438	97.0654	389.269
CPU (vectorized - no opt - 1 thread)	3.20216	13.6765	51.7357	208.745
CPU (basic - O2 - 1 thread)	1.44004	7.59237	28.6857	86.6324
CPU (vectorized - O2 - 1 thread)	0.844237	3.37243	13.7614	54.3744
CPU (basic - O3 - 1 thread)	0.528666	2.11186	8.45428	33.8729
CPU (vectorized - O3 - 1 thread)	0.850539	3.3954	13.8491	55.3561
CPU (basic - O3 - 4 threads)	0.529679	2.12084	8.43355	33.806
CPU (vectorized - O3 - 4 threads)	0.846091	3.37604	13.4985	54.034

Table 3: Runtimes of CPU and GPU convolution implementations

runtimes by up to 4× as work is split across CPU cores. Compared with the CPU, the GPU is faster overall due to its massive parallelism and memory hierarchy tuned for convolution-style workloads—effectively like running on a CPU with far more threads.

4.0 Evaluations

Device	Model	Process Tech.	Die Size (mm ²)	TDP (W)
CPU	Intel core i7-11700 (8 cores)	14nm	276	65
GPU	Nvidia GeForce RTX3070	8nm	393	220
FPGA	Arria 10 GX 1150	20nm	~400	70

Table 4: Specification of the three compute platforms

	Throughput (GOPS)	Power (W)	Energy Efficiency (GOPS/W)	Area Efficiency (GOPS/mm ²)
FPGA (20nm)	1.2510×10^3	11.98	104.42	3.5743
GPU (28nm)	15.786	165	0.095673	0.39864
GPU (scaled 20nm)	22.1	247.5	0.089293	0.055808
CPU (22nm)	163.72	84	1.9490	0.92497

Table 5: Device evaluations

Base on the results, FPGA performs way better than CPU and GPU. This is mostly caused by FPGA is only performing the calculation that directly contribute to the convolutions. CPU and GPU involves with lots of other operations that handles with the control flow and other setting which are required for the convolution.

Additionally, the CPU and GPU power figures reflect theoretical upper limits and are unlikely to match the actual draw during our runs. In contrast, the FPGA power numbers are more representative, coming from a power analysis that uses realistic toggle rates derived from a ModelSim simulation.

Furthermore, the area metrics also amplify the FPGA's advantage: compared with the CPU and GPU, a far larger fraction of the FPGA's counted area is actively dedicated to the convolution itself, rather than to general-purpose overhead needed by CPUs/GPUs.

In conclusion, the FPGA throughput and power figures are likely optimistic, since they are extrapolated from a single module. Scaling to near-full device utilization would not be linear, routing congestion and suboptimal placement can raise power and push down achievable clock frequency.

5.0 Appendix

Appendix A: FPGA Correctness

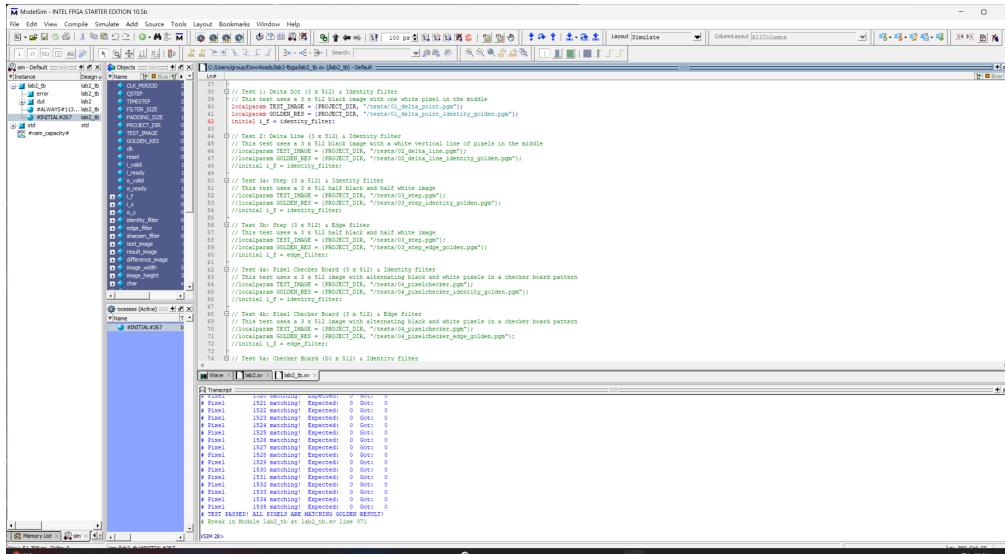


Figure 4: FPGA test 1 test bench

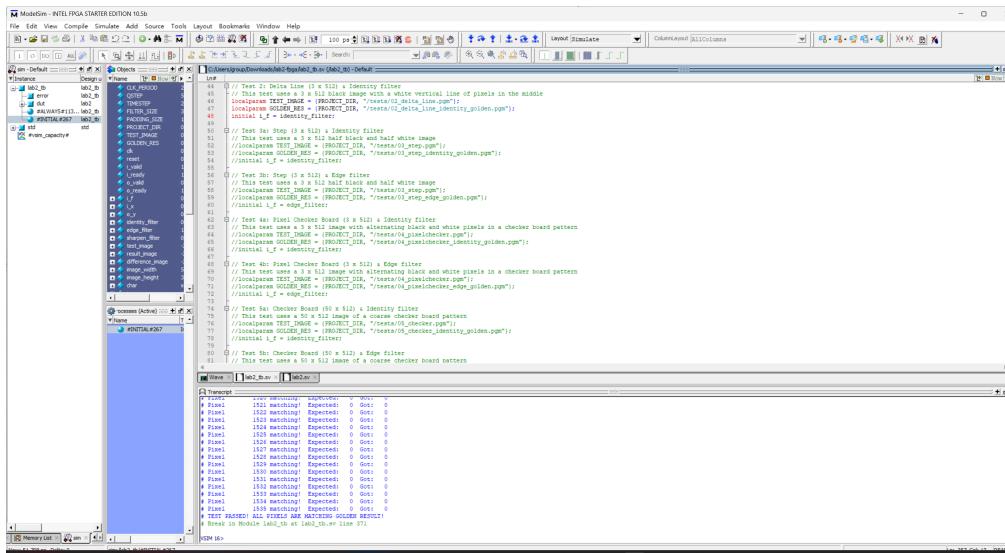


Figure 5: FPGA test 2 test bench

Assignment 2

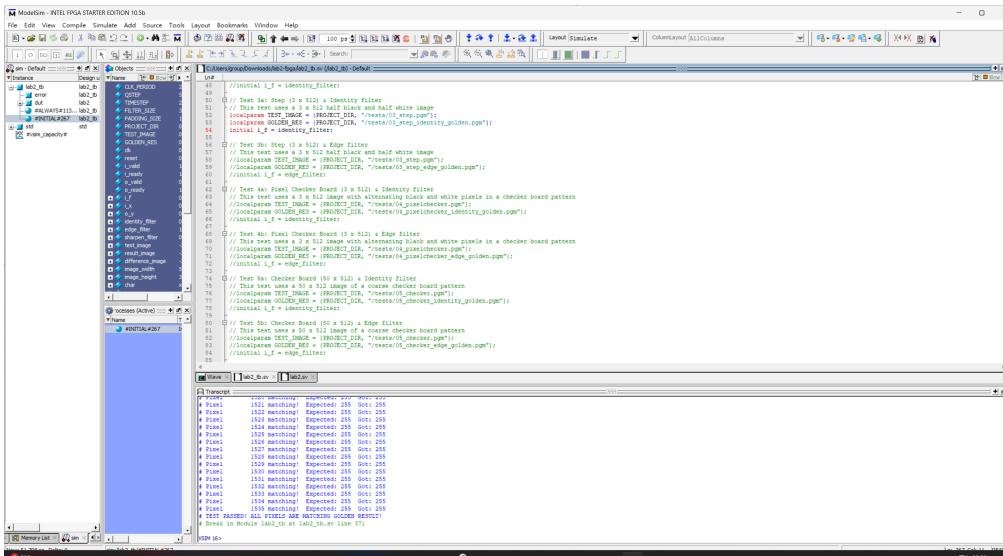


Figure 6: FPGA test 3a test bench

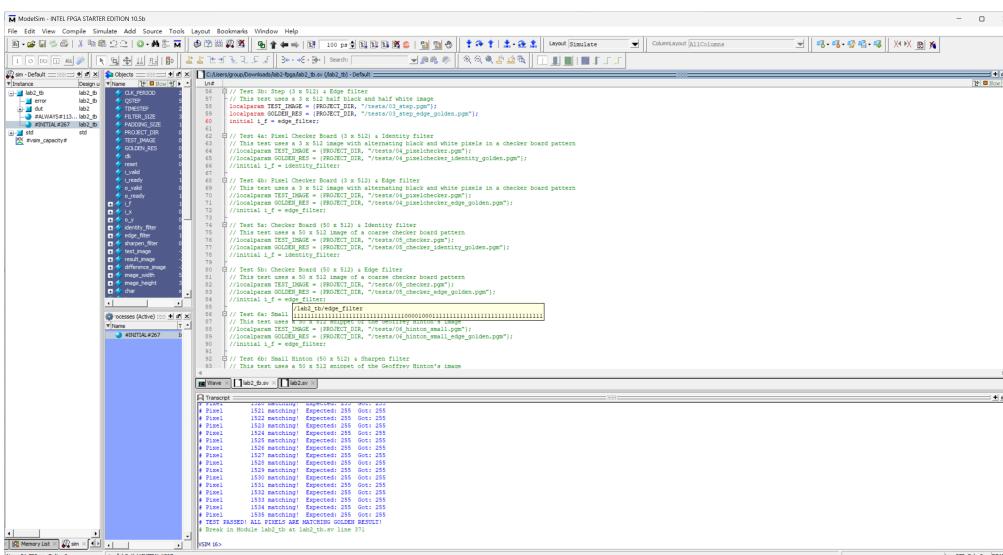


Figure 7: FPGA test 3b test bench

Assignment 2

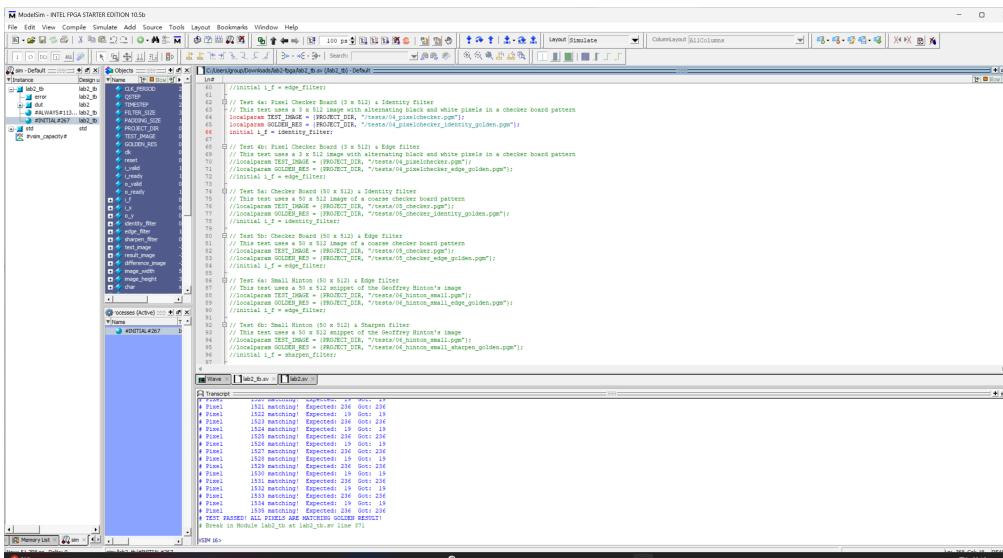


Figure 8: FPGA test 4a test bench

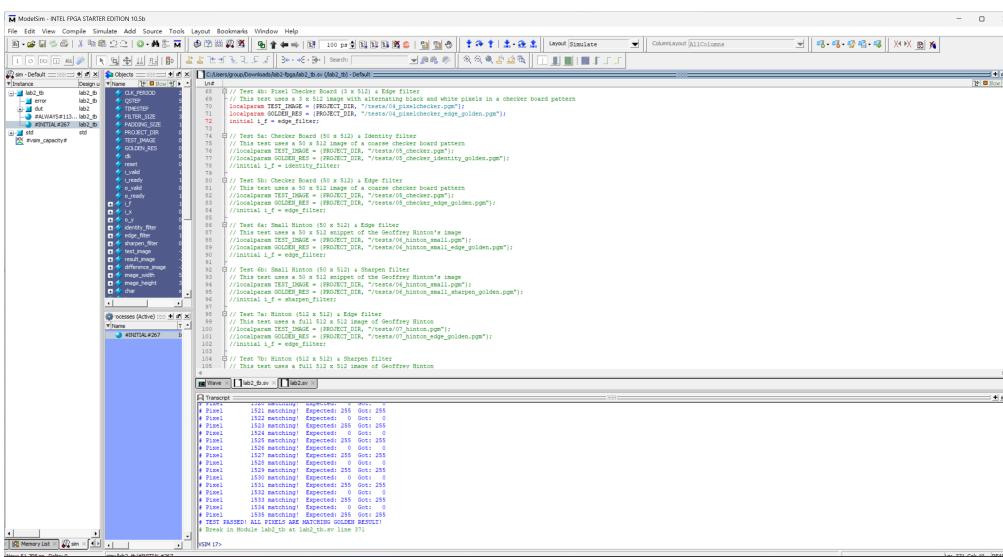


Figure 9: FPGA test 4b test bench

Assignment 2

ECE 1756
November 04, 2025

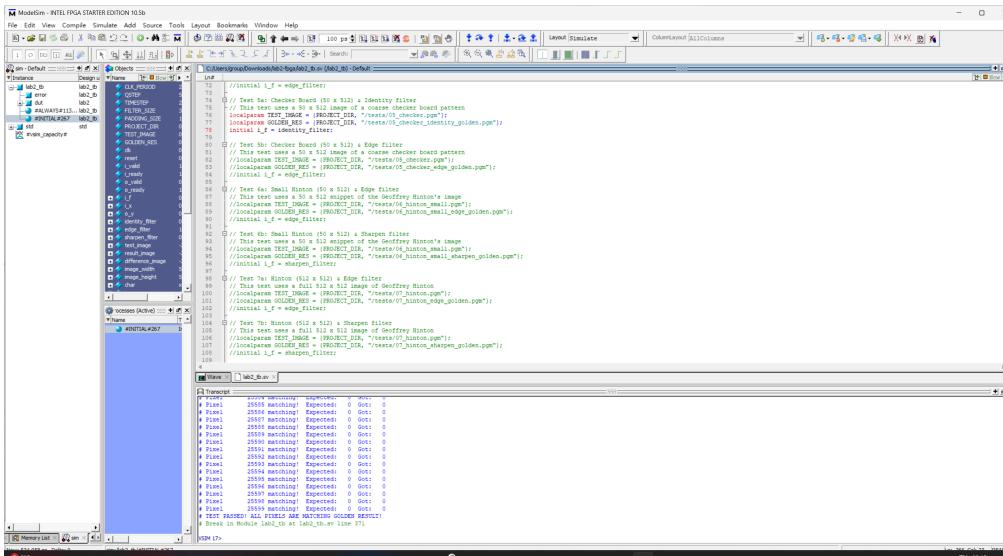


Figure 10: FPGA test 5a test bench

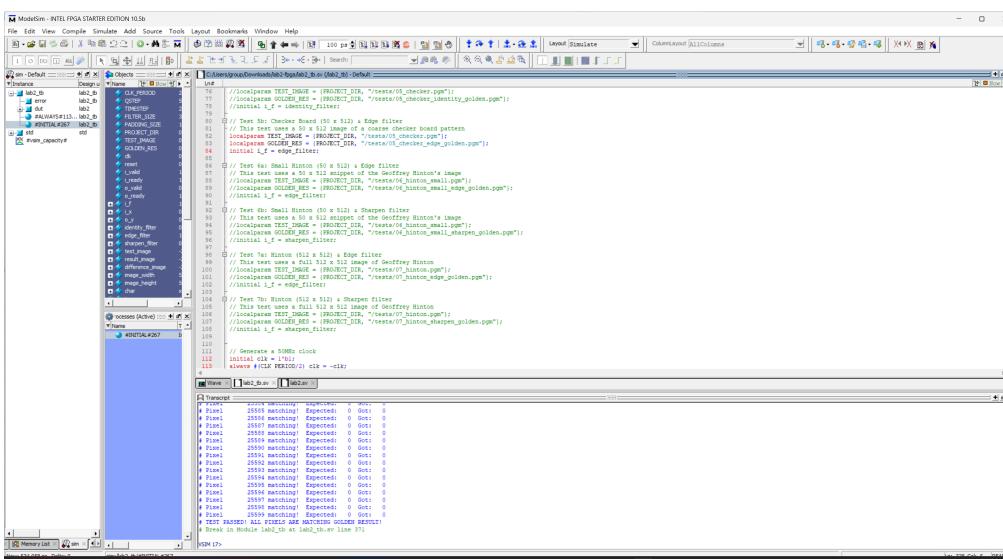


Figure 11: FPGA test 5b test bench

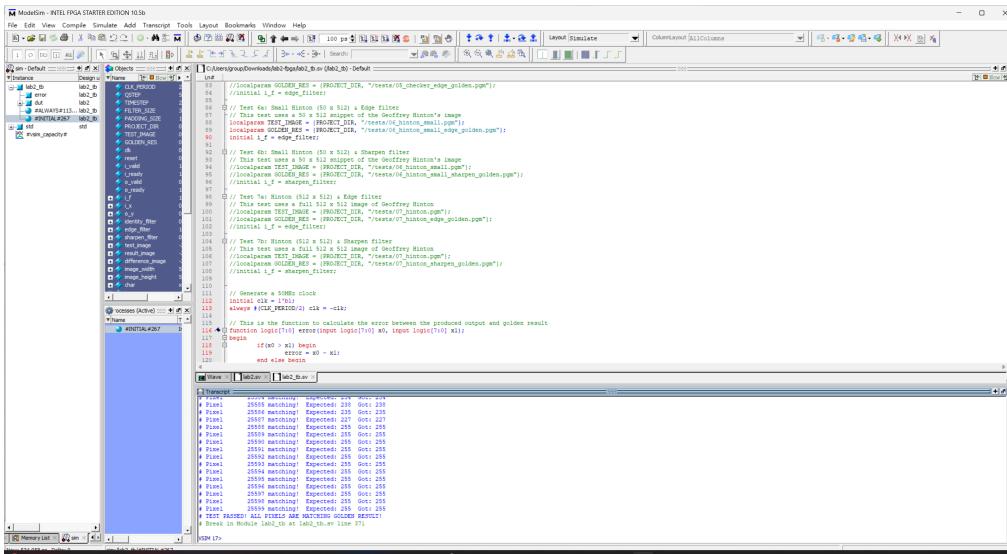


Figure 12: FPGA test 6a test bench

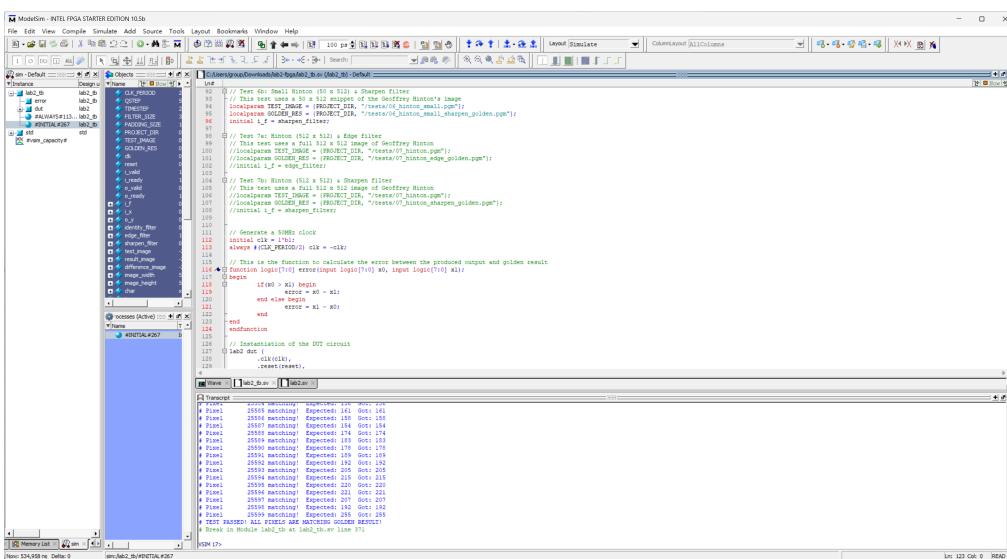


Figure 13: FPGA test 6b test bench

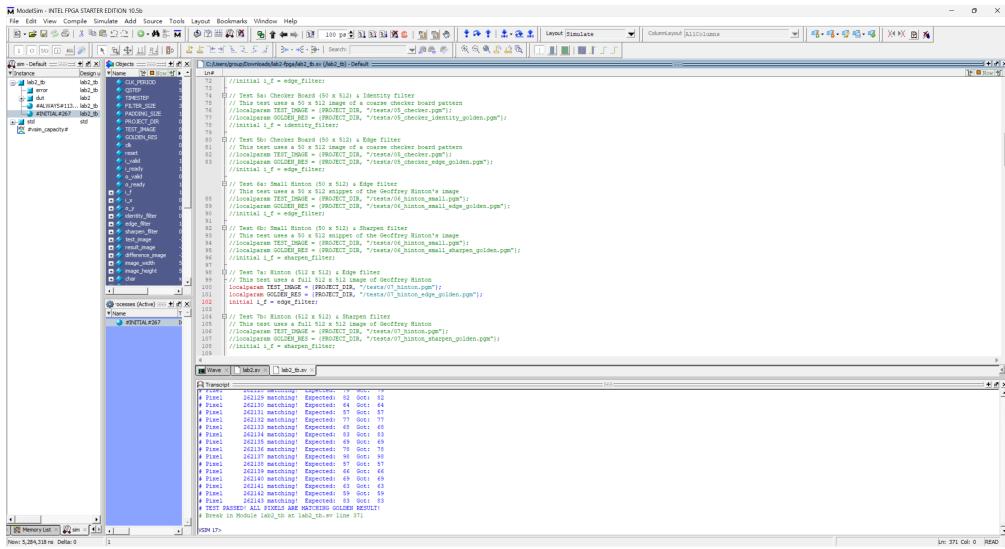


Figure 14: FPGA test 7a test bench

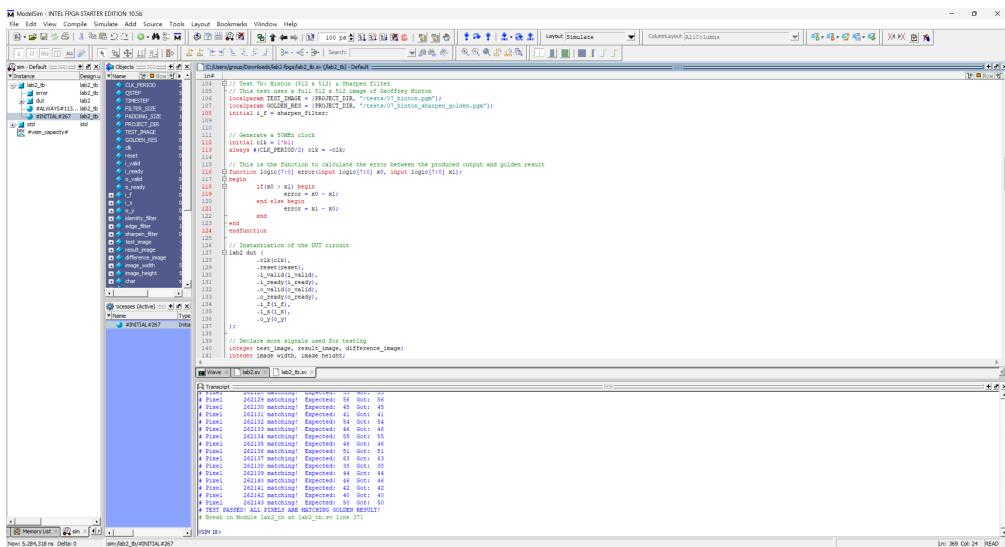


Figure 15: FPGA test 7b test bench

Appendix B: FPGA Implementation Results Calculations

ALM, DSP, BRAM Utilization

Flow Status	Successful - Tue Nov 04 17:40:45 2025
Quartus Prime Version	20.1.0 Build 711 06/05/2020 SJ Standard Edition
Revision Name	lab2
Top-level Entity Name	lab2
Family	Arria 10
Device	10AX115N1F45I1SG
Timing Models	Final
Logic utilization (in ALMs)	2,220 / 427,200 (< 1 %)
Total pins	1 / 992 (< 1 %)
Total virtual pins	93
Total block memory bits	0 / 55,562,240 (0 %)
Total DSP Blocks	9 / 1,518 (< 1 %)
Total HSSI RX channels	0 / 48 (0 %)
Total HSSI TX channels	0 / 48 (0 %)
Total PLLs	0 / 112 (0 %)

Figure 16: FPGA ALM, DSP and BRAM utilization

Maximum Operating Frequency

```

Analyzing Slow 900mV 100C Model
> A 332148 Timing requirements not met
> i 332146 Worst-case setup slack is -1.078
> i 332146 Worst-case hold slack is 0.045
i 332140 No Recovery paths to report
i 332140 No Removal paths to report
> i 332146 Worst-case minimum pulse width slack is -1.150
Analyzing Slow 900mV OC Model
> i 332146 Worst-case setup slack is -1.261
> i 332146 Worst-case hold slack is 0.045
i 332140 No Recovery paths to report

```

Figure 17: FPGA setup and hold times

The screenshot shows the Xilinx Vivado IDE interface. The top menu bar includes 'lab2.sv', 'Compilation Report - lab2', and 'lab2_tb.sv'. The left sidebar has a 'Table of Contents' section with various flow steps like Flow Summary, Flow Settings, etc., and a 'Timing Analyzer' section with sub-options like Summary, Parallel Compilation, SDC File List, Clocks, Slow 900mV 100C Mode, and Slow 900mV OC Model. The main panel displays a table titled 'Slow 900mV OC Model Fmax Summary' with one row of data:

	Fmax	Restricted Fmax	Clock Name	Note
1	442.28 MHz	442.28 MHz	clk	

Figure 18: FPGA maximum operating frequency

Cycles for Test 7a (Hinton)

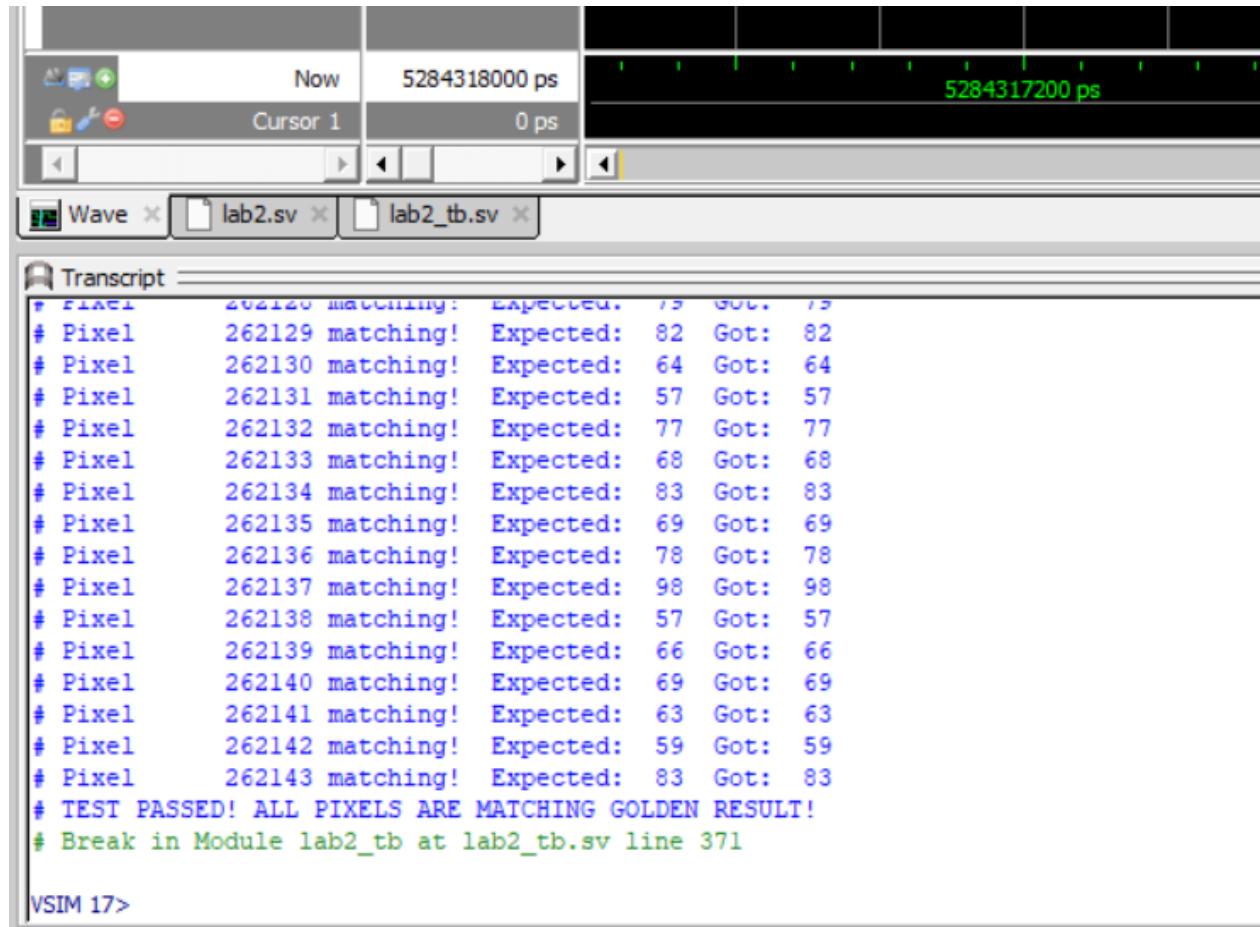


Figure 19: Test 7a total elapsed time

$$\text{No. of cycles} = 5284318000 / 20000 = 264216 \text{ cycles}$$

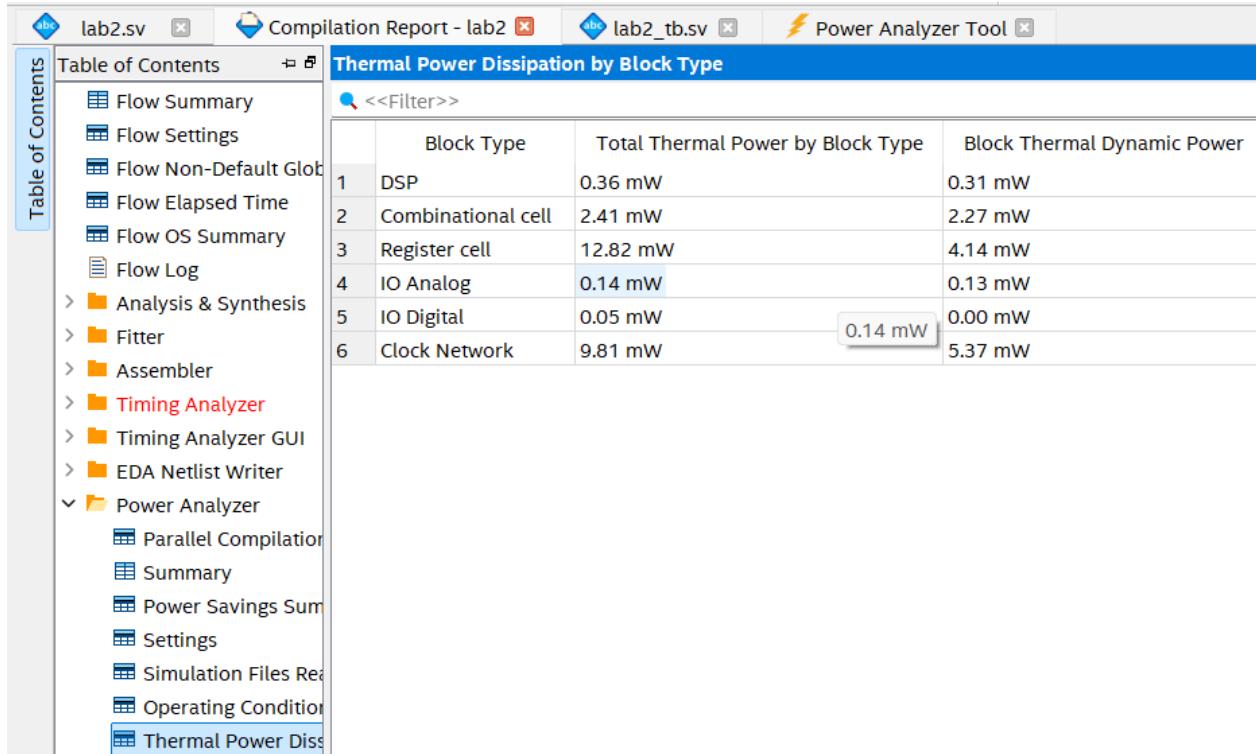
Dynamic Power for one module @ 442.28MHz

Figure 20: Power dissipation by block type

$$\text{Total Dynamic Power @}50\text{MHz} = 0.31 + 2.27 + 4.14 + 0.13 + 5.37 = 12.22 \text{ mW}$$

$$\text{Total Dynamic Power @}441.5\text{MHz} = 12.22 * (442.28/50) = 108.09 \text{ mW} = 108.09 \times 10^{-3} \text{ W}$$

Throughput of one module (GOPS)

To compute the convolution of one pixel, 9 multiplications and a total of 8 additions needs to be performed over a period of 5 cycles (output ready by the 6th cycle). For a 512x512 picture as is given in test 7a, that would take a total of:

$$(512 * 512) * (9 + 8) = 4456448 \text{ operations}$$

With a frequency of 442.28 MHz, 264216 cycles is equal to:

$$(1/441.5 \times 10^6) * 264216 \approx 5.9845 \times 10^{-4} \text{ seconds}$$

Which give us a total of:

$$4456448/5.9845 \times 10^{-4} \approx 7.4466 \times 10^9 = 7.4466 \text{ GOPs}$$

Throughput of a full device (GOPS)

$$\text{Max. copies (based on ALM count)} = 427200/2220 \approx 192$$

$$\text{Max. copies (based on DSP count)} = 1518/9 \approx 168$$

Since DSP count is the limiting factor, approximate maximum number of copies per device is 168. As a result, the throughput of a full device:

$$7.4466 \times 10^9 * 168 \approx 1.2510 \times 10^{12} = 1.2510 \times 10^3 \text{ GOPs}$$

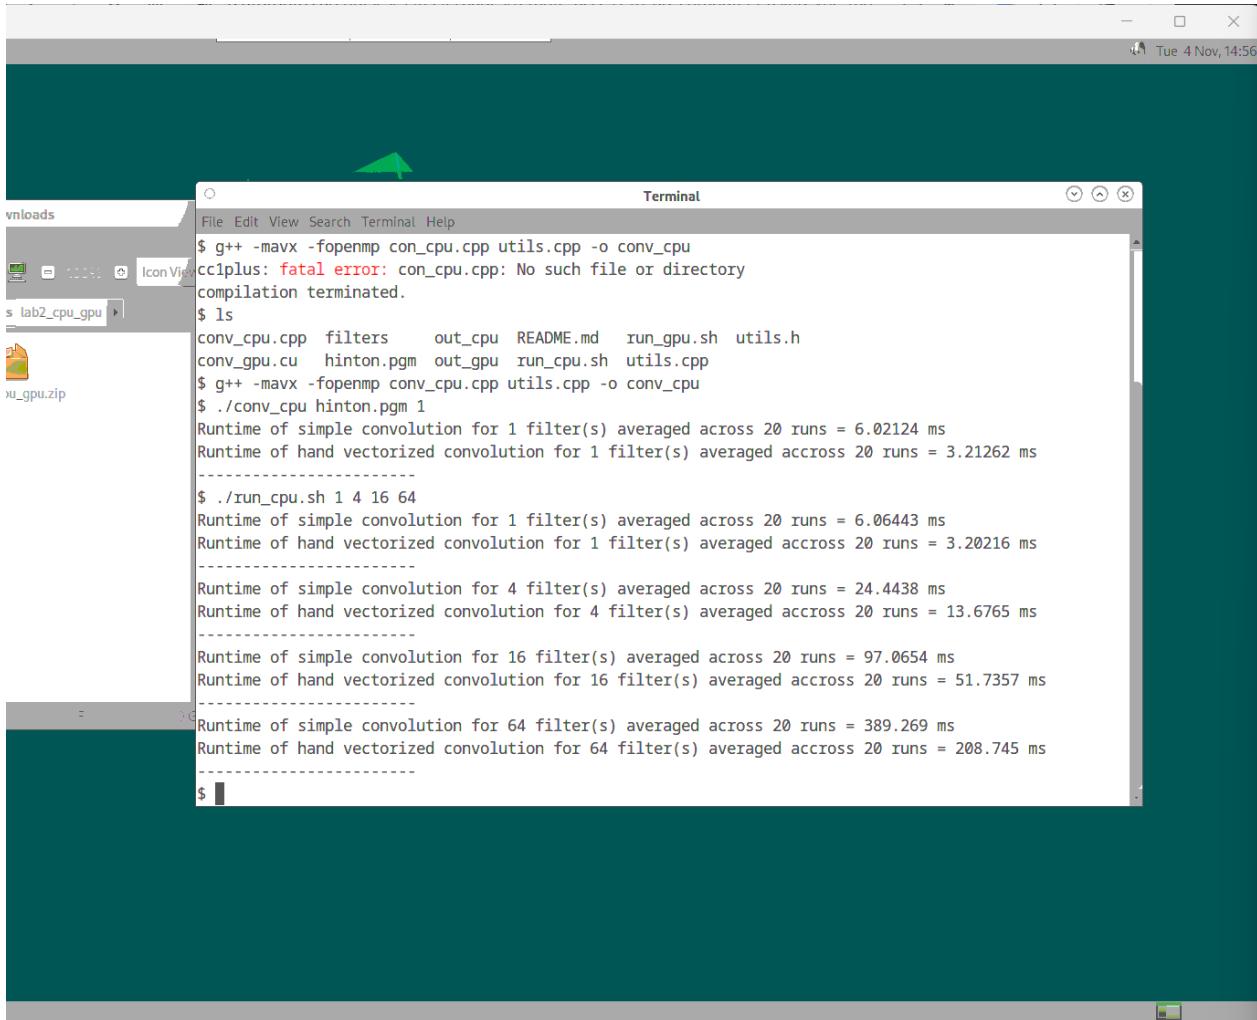
Total Power for a full device (W)

Assuming a fully packed device with 168 modules as calculated before (note that IO and clock dynamic power are not multiplied):

$$\text{Total Dynamic Power @50MHz} = (0.31 + 2.27 + 4.14) * 168 + 0.13 + 5.37 = 1134.46 \text{ mW}$$

$$\text{Total Dynamic Power @442.28MHz} = 1134.46 * (442.28/50) = 10034.98 \text{ mW}$$

$$\begin{aligned} \text{Total Power @442.28MHz} &= 10034.98 + (0.36 + 2.41 + 12.82) * 168 + 0.14 + 0.05 + 9.81 \\ &= 12664.1 \text{ mW} \approx 12.66 \text{ W} \end{aligned}$$

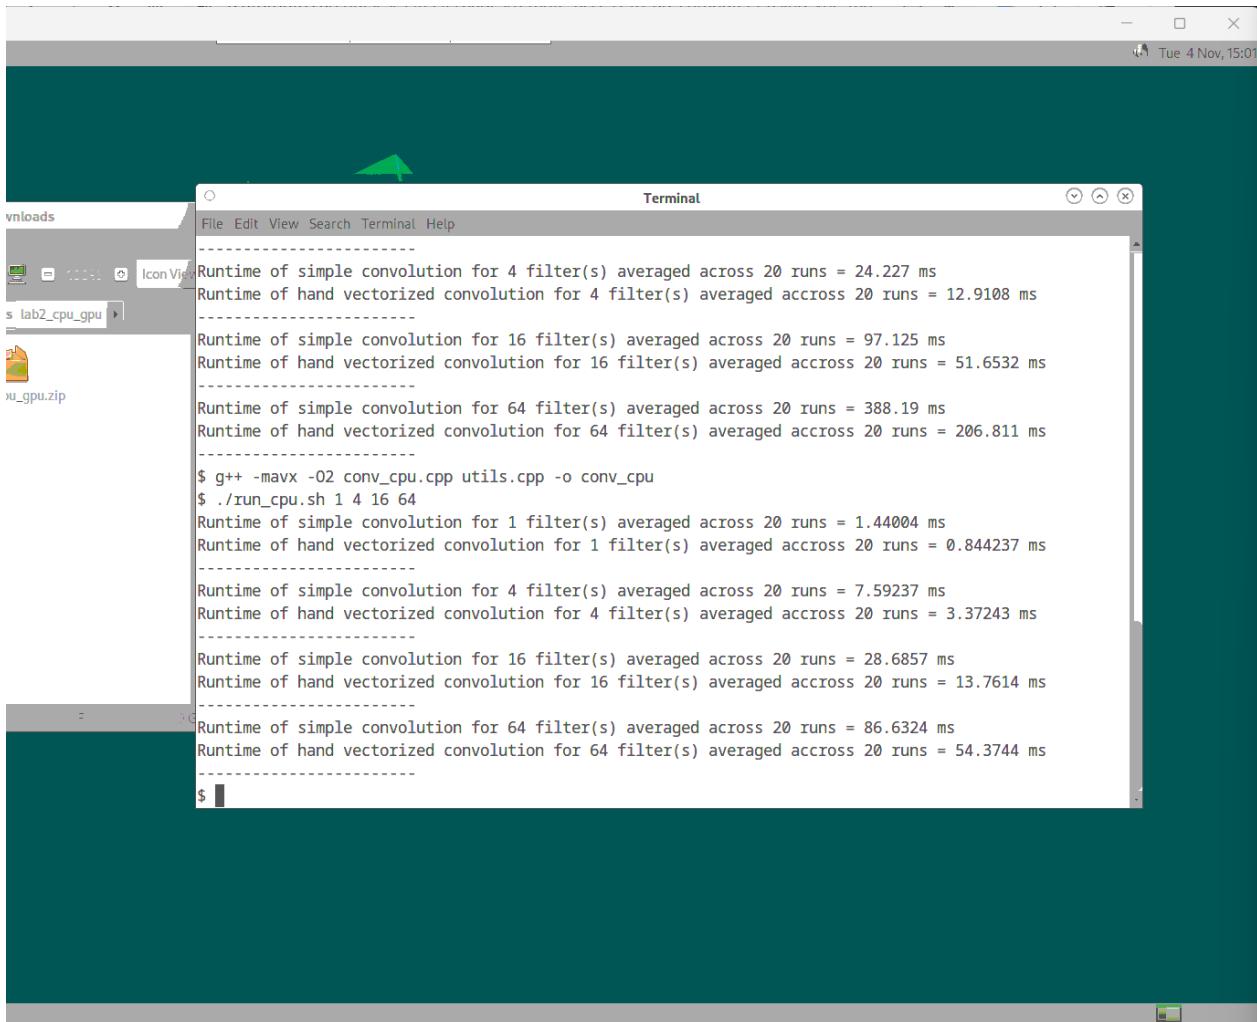
Appendix C: Runtimes for CPU and GPU implementations

The screenshot shows a Linux desktop environment with a terminal window open. The terminal window title is "Terminal". The terminal content displays the following command-line session:

```
$ g++ -fopenmp conv_cpu.cpp utils.cpp -o conv_cpu
cc1plus: fatal error: conv_cpu.cpp: No such file or directory
compilation terminated.

$ ls
conv_cpu.cpp filters out_cpu README.md run_gpu.sh utils.h
conv_gpu.cu hinton.pgm out_gpu run_cpu.sh utils.cpp
$ g++ -fopenmp conv_cpu.cpp utils.cpp -o conv_cpu
$ ./conv_cpu hinton.pgm 1
Runtime of simple convolution for 1 filter(s) averaged across 20 runs = 6.02124 ms
Runtime of hand vectorized convolution for 1 filter(s) averaged accross 20 runs = 3.21262 ms
-----
$ ./run_cpu.sh 1 4 16 64
Runtime of simple convolution for 1 filter(s) averaged across 20 runs = 6.06443 ms
Runtime of hand vectorized convolution for 1 filter(s) averaged accross 20 runs = 3.20216 ms
-----
Runtime of simple convolution for 4 filter(s) averaged across 20 runs = 24.4438 ms
Runtime of hand vectorized convolution for 4 filter(s) averaged accross 20 runs = 13.6765 ms
-----
Runtime of simple convolution for 16 filter(s) averaged across 20 runs = 97.0654 ms
Runtime of hand vectorized convolution for 16 filter(s) averaged accross 20 runs = 51.7357 ms
-----
Runtime of simple convolution for 64 filter(s) averaged across 20 runs = 389.269 ms
Runtime of hand vectorized convolution for 64 filter(s) averaged accross 20 runs = 208.745 ms
-----
```

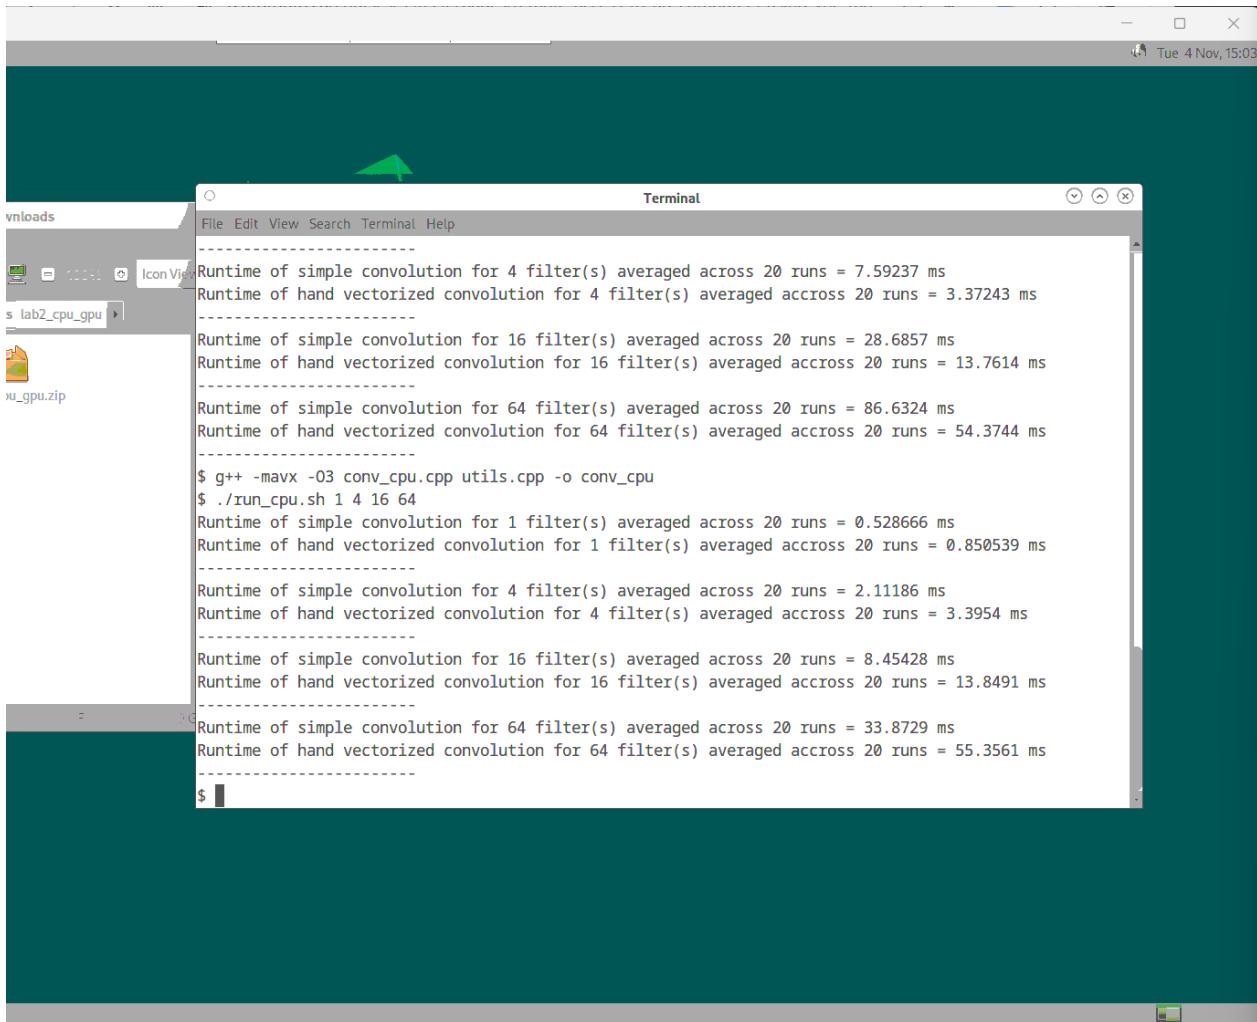
Figure 21: CPU no optimization runtimes



The screenshot shows a Linux desktop environment with a terminal window open. The terminal window title is "Terminal". The terminal displays the following output:

```
Runtime of simple convolution for 4 filter(s) averaged across 20 runs = 24.227 ms
Runtime of hand vectorized convolution for 4 filter(s) averaged accross 20 runs = 12.9108 ms
-----
Runtime of simple convolution for 16 filter(s) averaged across 20 runs = 97.125 ms
Runtime of hand vectorized convolution for 16 filter(s) averaged accross 20 runs = 51.6532 ms
-----
Runtime of simple convolution for 64 filter(s) averaged across 20 runs = 388.19 ms
Runtime of hand vectorized convolution for 64 filter(s) averaged accross 20 runs = 206.811 ms
-----
$ g++ -mavx -O2 conv_cpu.cpp utils.cpp -o conv_cpu
$ ./run_cpu.sh 1 4 16 64
Runtime of simple convolution for 1 filter(s) averaged across 20 runs = 1.44004 ms
Runtime of hand vectorized convolution for 1 filter(s) averaged accross 20 runs = 0.844237 ms
-----
Runtime of simple convolution for 4 filter(s) averaged across 20 runs = 7.59237 ms
Runtime of hand vectorized convolution for 4 filter(s) averaged accross 20 runs = 3.37243 ms
-----
Runtime of simple convolution for 16 filter(s) averaged across 20 runs = 28.6857 ms
Runtime of hand vectorized convolution for 16 filter(s) averaged accross 20 runs = 13.7614 ms
-----
Runtime of simple convolution for 64 filter(s) averaged across 20 runs = 86.6324 ms
Runtime of hand vectorized convolution for 64 filter(s) averaged accross 20 runs = 54.3744 ms
-----
```

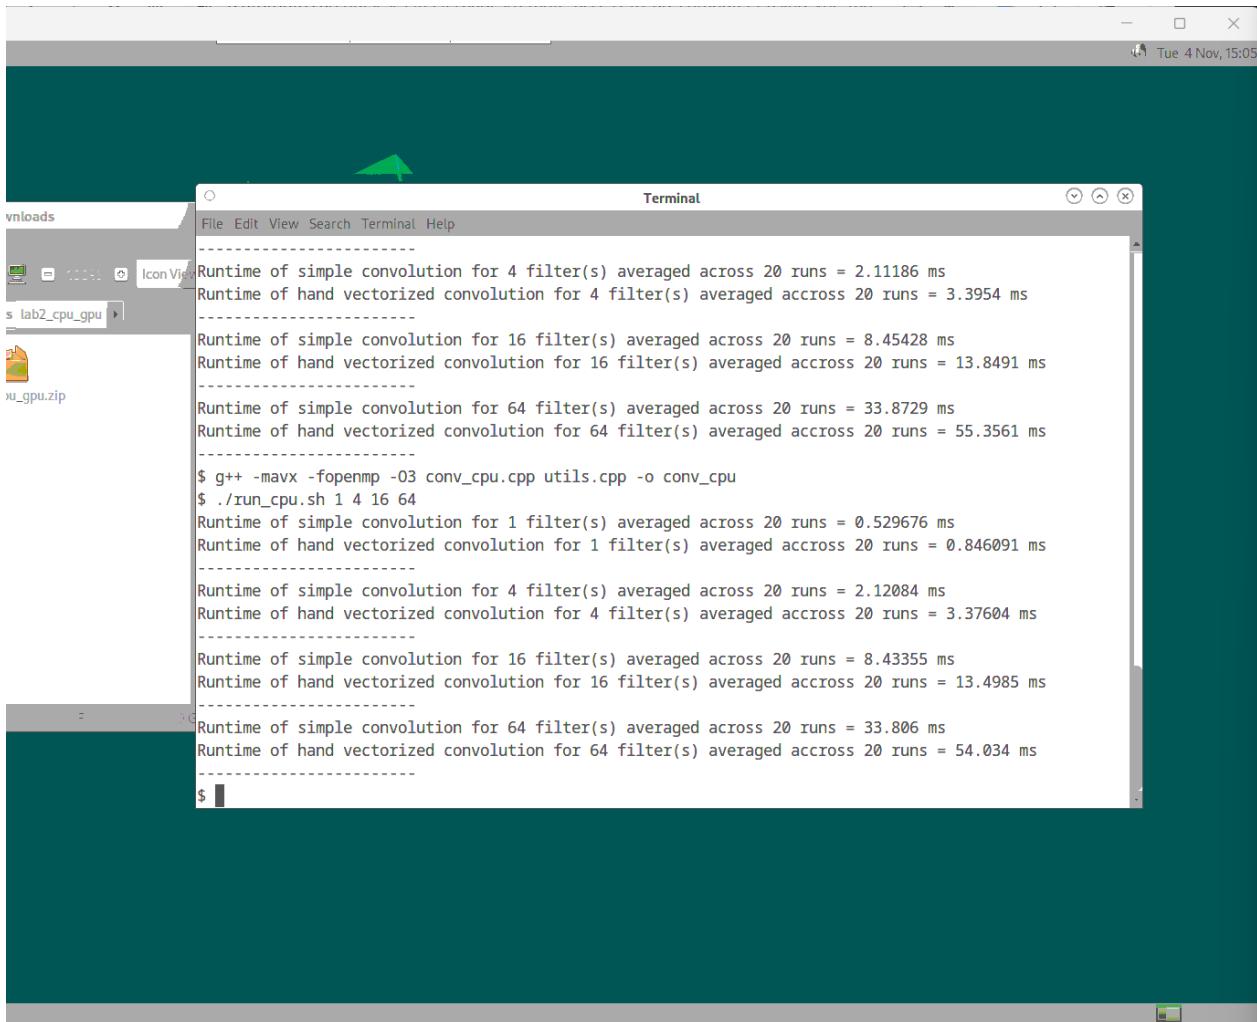
Figure 22: CPU runtimes with O2 optimization



The screenshot shows a Linux desktop environment with a terminal window open. The terminal window title is "Terminal". Inside the terminal, there is a series of runtime measurements for convolution operations. The measurements are grouped by filter count (4, 16, 64) and show two types of convolution: "simple convolution" and "hand vectorized convolution". The times are averaged across 20 runs. The terminal also shows the compilation command and the execution of a script to run the benchmarks.

```
Runtime of simple convolution for 4 filter(s) averaged across 20 runs = 7.59237 ms
Runtime of hand vectorized convolution for 4 filter(s) averaged accross 20 runs = 3.37243 ms
-----
Runtime of simple convolution for 16 filter(s) averaged across 20 runs = 28.6857 ms
Runtime of hand vectorized convolution for 16 filter(s) averaged accross 20 runs = 13.7614 ms
-----
Runtime of simple convolution for 64 filter(s) averaged across 20 runs = 86.6324 ms
Runtime of hand vectorized convolution for 64 filter(s) averaged accross 20 runs = 54.3744 ms
-----
$ g++ -O3 conv_cpu.cpp utils.cpp -o conv_cpu
$ ./run_cpu.sh 1 4 16 64
Runtime of simple convolution for 1 filter(s) averaged across 20 runs = 0.528666 ms
Runtime of hand vectorized convolution for 1 filter(s) averaged accross 20 runs = 0.850539 ms
-----
Runtime of simple convolution for 4 filter(s) averaged across 20 runs = 2.11186 ms
Runtime of hand vectorized convolution for 4 filter(s) averaged accross 20 runs = 3.3954 ms
-----
Runtime of simple convolution for 16 filter(s) averaged across 20 runs = 8.45428 ms
Runtime of hand vectorized convolution for 16 filter(s) averaged accross 20 runs = 13.8491 ms
-----
Runtime of simple convolution for 64 filter(s) averaged across 20 runs = 33.8729 ms
Runtime of hand vectorized convolution for 64 filter(s) averaged accross 20 runs = 55.3561 ms
-----
```

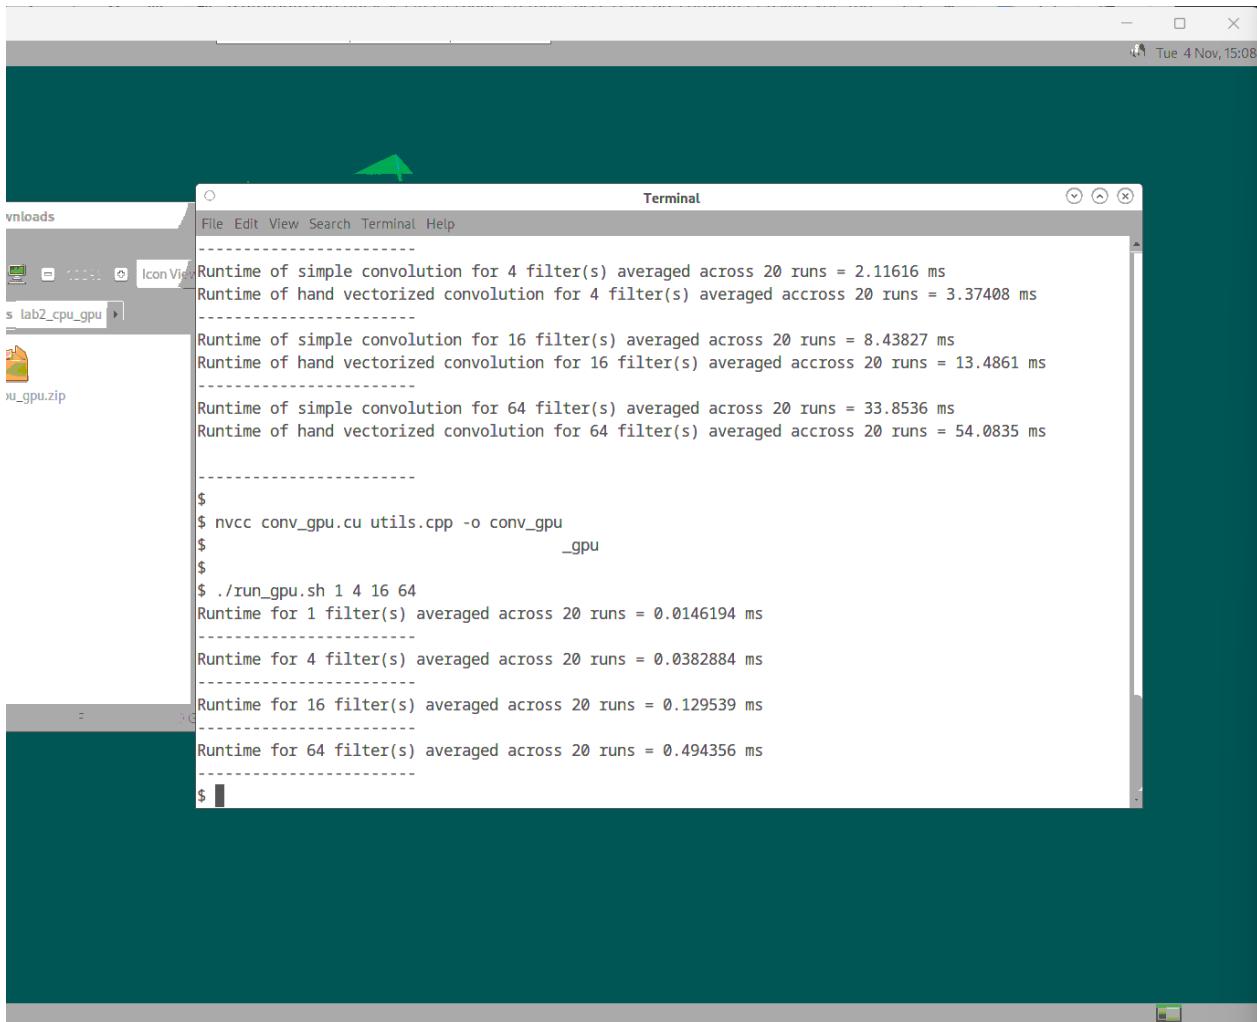
Figure 23: CPU runtimes with O3 optimization



The screenshot shows a Linux desktop environment with a terminal window open. The terminal window title is "Terminal". Inside the terminal, the user has run a script to measure convolution times. The output shows results for simple and hand-vectorized convolutions with 4, 16, and 64 filters, comparing them across 20 runs. The user then compiles a program named "conv_cpu" and runs it with different filter counts. The terminal also shows the compilation command and the execution of the script.

```
vnloads 100% 0:00:00.000000 100% 0:00:00.000000
Icon View
File Edit View Search Terminal Help
-----
Runtime of simple convolution for 4 filter(s) averaged across 20 runs = 2.11186 ms
Runtime of hand vectorized convolution for 4 filter(s) averaged accross 20 runs = 3.3954 ms
-----
Runtime of simple convolution for 16 filter(s) averaged across 20 runs = 8.45428 ms
Runtime of hand vectorized convolution for 16 filter(s) averaged accross 20 runs = 13.8491 ms
-----
Runtime of simple convolution for 64 filter(s) averaged across 20 runs = 33.8729 ms
Runtime of hand vectorized convolution for 64 filter(s) averaged accross 20 runs = 55.3561 ms
-----
$ g++ -fopenmp -O3 conv_cpu.cpp utils.cpp -o conv_cpu
$ ./run_cpu.sh 1 4 16 64
Runtime of simple convolution for 1 filter(s) averaged across 20 runs = 0.529676 ms
Runtime of hand vectorized convolution for 1 filter(s) averaged accross 20 runs = 0.846091 ms
-----
Runtime of simple convolution for 4 filter(s) averaged across 20 runs = 2.12084 ms
Runtime of hand vectorized convolution for 4 filter(s) averaged accross 20 runs = 3.37604 ms
-----
Runtime of simple convolution for 16 filter(s) averaged across 20 runs = 8.43355 ms
Runtime of hand vectorized convolution for 16 filter(s) averaged accross 20 runs = 13.4985 ms
-----
Runtime of simple convolution for 64 filter(s) averaged across 20 runs = 33.806 ms
Runtime of hand vectorized convolution for 64 filter(s) averaged accross 20 runs = 54.034 ms
-----
$
```

Figure 24: CPU runtimes with O3 optimization, multithreaded



The screenshot shows a Linux desktop environment with a terminal window open. The terminal window title is "Terminal". The terminal displays the following output:

```
Runtime of simple convolution for 4 filter(s) averaged across 20 runs = 2.11616 ms
Runtime of hand vectorized convolution for 4 filter(s) averaged accross 20 runs = 3.37408 ms
-----
Runtime of simple convolution for 16 filter(s) averaged across 20 runs = 8.43827 ms
Runtime of hand vectorized convolution for 16 filter(s) averaged accross 20 runs = 13.4861 ms
-----
Runtime of simple convolution for 64 filter(s) averaged across 20 runs = 33.8536 ms
Runtime of hand vectorized convolution for 64 filter(s) averaged accross 20 runs = 54.0835 ms
-----
$ 
$ nvcc conv_gpu.cu utils.cpp -o conv_gpu
$                               _gpu
$ 
$ ./run_gpu.sh 1 4 16 64
Runtime for 1 filter(s) averaged across 20 runs = 0.0146194 ms
-----
Runtime for 4 filter(s) averaged across 20 runs = 0.0382884 ms
-----
Runtime for 16 filter(s) averaged across 20 runs = 0.129539 ms
-----
Runtime for 64 filter(s) averaged across 20 runs = 0.494356 ms
-----
```

Figure 25: GPU runtimes

Appendix D: Evaluation Calculations

FPGA related

$$\text{Energy Efficiency} = 1.2510 \times 10^3 / 11.98 \approx 104.42 \text{ GOPS/W}$$

$$\text{Area Efficiency} = 1.2510 \times 10^3 / 400 \approx 3.1275 \text{ GOPS/mm}^2$$

CPU related

```
/*
 * This function implements the 2D convolution simply using nested for loops assuming that
 * the input image is already padded. It takes as an input pointers for input and output
 * images as well as convolution filter coefficients and number of filters used.
 */
void cpu_conv2d(float* input_image, float* filter, float* output_image, int image_width, int image_height, int num_filters){
    // Declare variables used throughout the function
    float accum;
    int xx, yy;

    // Calculate the padded image size
    int padded_image_width = image_width + (2*FILTER_RADIUS);

    for(int filter_id = 0; filter_id < num_filters; filter_id++){ // For each filter
        for(int image_y = 0; image_y < image_height; image_y++){ // For each y location of the output image
            for(int image_x = 0; image_x < image_width; image_x++){ // For each x location of the output image
                // Start a new accumulation result
                accum = 0;

                // Compute the dot-product of the filter and the input image
                for(int filter_y = 0; filter_y < FILTER_SIZE; filter_y++){
                    for(int filter_x = 0; filter_x < FILTER_SIZE; filter_x++){
                        xx = image_x + filter_x;
                        yy = image_y + filter_y;
                        accum += input_image[(yy * padded_image_width) + xx] * filter[(filter_id * FILTER_SIZE * FILTER_SIZE) + (filter_y * FILTER_SIZE) + filter_x];
                    }
                }

                // Store the convolution result in the corresponding pixel location
                output_image[filter_id * image_width * image_height + (image_y * image_width) + image_x] = accum;
            }
        }
    }
}
```

Figure 26: CPU basic code

Based on the figure above, we can estimate the number of operations that are performed for the convolution in the basic implementation on the CPU. Number of operations within the most inner loop = 11. FILTER_SIZE = 3 image.height = image.width = 512 + 2 num_filters = 64 Margin of error = 10%

$$\begin{aligned} \text{Total Operations} &= (((11 * 3 * 3) + 5) * 514 * 514 * 64) * 1.10 / 11.8149 \times 10^{-3} \\ &\approx 163.72 \text{ GOPS} \end{aligned}$$

$$\begin{aligned} \text{Energy Efficiency} &= 163.72 / 65 \approx 2.518 \text{ GOPS/W} \\ \text{Area Efficiency} &= 163.72 / 276 \approx 0.593 \text{ GOPS/mm}^2 \end{aligned}$$

GPU related

```

global void gpu conv2d(int* input image, int* output image, int image width, int image height) {
    // Allocate space in shared memory to load an image tile
    shared int image tile[TILE LOAD SIZE][TILE LOAD SIZE];

    // Declare variables to use throughout the kernel
    int thread id, tile location x, tile location y;
    int image location x, image location y, pixel id;

    // Load tiles from Global memory to Shared memory for faster access
    for (int itr = 0; itr <= (TILE LOAD SIZE * TILE LOAD SIZE) / (TILE SIZE * TILE SIZE); itr++){
        // Calculate destination x and y indecies
        thread id = (threadIdx.y * TILE SIZE) + threadIdx.x + (itr * TILE SIZE * TILE SIZE);
        tile location y = thread id / TILE LOAD SIZE;
        tile location x = thread id % TILE LOAD SIZE;

        // Calculate source pixel index
        image location y = blockIdx.y * TILE SIZE + tile location y - FILTER RADIUS;
        image location x = blockIdx.x * TILE SIZE + tile location x - FILTER RADIUS;
        pixel id = (image location y * image width) + image location x;

        // Load pixels
        if (tile location y < TILE LOAD SIZE && tile location x < TILE LOAD SIZE){
            if (image location y >= 0 && image location y < image height
                && image location x >= 0 && image location x < image width){
                    image tile[tile location y][tile location x] = input image[pixel id];
                } else {
                    image tile[tile location y][tile location x] = 0;
                }
        }
    }
    syncthreads();
}

```

Figure 27: GPU code

Based on the figure above, we can estimate the number of operations that are performed for the convolution in the implementation on the GPU. Number of operations within loop = 15
 $TILE_SIZE = 32$ $TILE_LOAD_SIZE = 32 + 3 - 1 = 34$

```

// Perform the 2D convolution
int accum = 0;
int y, x, z;
z = blockIdx.z;
for (y = 0; y < FILTER SIZE; y++) {
    for (x = 0; x < FILTER SIZE; x++) {
        accum += image tile[threadIdx.y + y][threadIdx.x + x] * device filter[(z * FILTER SIZE * FILTER SIZE) + (y * FILTER SIZE) + x];
    }
}

// Write the output
y = blockIdx.y * TILE SIZE + threadIdx.y;
x = blockIdx.x * TILE SIZE + threadIdx.x;
if (y < image height && x < image width){
    output image[(z * image width * image height) + (y * image width) + x] = accum;
}
syncthreads();

```

Figure 28: GPU code

Number of operations within loop = 9 $FILTER_SIZE = 3$ Above repeated for each pixel.
Number of operations outside loop = 9 Margin of error = 10%

$$\begin{aligned} \text{Total Operations} &= (15*((34*34)/(32*32)) + (9*3*3)*512*512 + 9) * 1.10 / 1.47965 \times 10^{-3} \\ &\approx 15.786 \text{ GOPS} \end{aligned}$$

$$\text{Energy Efficiency} = 15.786 / 165 \approx 0.095673 \text{ GOPS/W}$$

$$\text{Area Efficiency} = 15.786 / 393 \approx 0.04016 \text{ GOPS/mm}^2$$

GPU Estimation

$$\text{Throughput} = 15.786 * (8/20) \approx 6.31 \text{ GOPS}$$

$$\text{Power} = 165 * 1.5 \approx 247.5 \text{ W}$$

$$\text{Energy Efficiency} = 22.1 / 247.5 \approx 0.089293 \text{ GOPS/W}$$

$$\text{Area Efficiency} = 22.1 / 393 \approx 0.05623 \text{ GOPS/mm}^2$$

Appendix E: .sv file for FPGA Implementation

```

1 // This module implements 2D convolution between a 3x3 filter and a
2   ↵ 512-pixel-wide image of any height.
3 // It is assumed that the input image is padded with zeros such that the
4   ↵ input and output images have
5 // the same size. The filter coefficients are symmetric in the x-direction
6   ↵ (i.e. f[0][0] = f[0][2],
7 // f[1][0] = f[1][2], f[2][0] = f[2][2] for any filter f) and their values
8   ↵ are limited to integers
9 // (but can still be positive or negative). The input image is grayscale
10  ↵ with 8-bit pixel values ranging
11 // from 0 (black) to 255 (white).
12 module lab2 (
13     input  clk,                      // Operating clock
14     input  reset,                    // Active-high reset signal
15     ↵  (reset when set to 1)
16     input [71:0] i_f,               // Nine 8-bit signed convolution
17       ↵ filter coefficients in row-major format (i.e. i_f[7:0] is
18       ↵ f[0][0], i_f[15:8] is f[0][1], etc.)
19     input  i_valid,                // Set to 1 if input pixel is
20       ↵ valid
21     input  i_ready,                // Set to 1 if consumer block
22       ↵ is ready to receive a new pixel
23     input [7:0] i_x,                 // Input pixel value (8-bit
24       ↵ unsigned value between 0 and 255)
25     output o_valid,               // Set to 1 if output pixel
26       ↵ is valid
27     output o_ready,                // Set to 1 if this block is
28       ↵ ready to receive a new pixel
29     output [7:0] o_y,              // Output pixel value (8-bit
30       ↵ unsigned value between 0 and 255)
31 );
32
33 localparam FILTER_SIZE = 3;        // Convolution filter dimension (i.e.
34   ↵ 3x3)
35 localparam PIXEL_DATAW = 8;        // Bit width of image pixels and filter
36   ↵ coefficients (i.e. 8 bits)localparam int FILTER_SIZE = 3;
37
38 // The following code is intended to show you an example of how to use
39   ↵ parameters and

```

```
23 // for loops in SystemVerilog. It also arranges the input filter coefficients
24 // into a nicely-arranged and easy-to-use 2D array of registers. However,
25 // you can ignore
26 // this code and not use it if you wish to.
27
28 logic signed [PIXEL_DATAW-1:0] r_f [FILTER_SIZE-1:0][FILTER_SIZE-1:0]; // 2D
29 // array of registers for filter coefficients
30 integer col, row, buf_i, i; // variables to use in the for loop
31
32 /*always_ff @ (posedge clk) begin
33     // If reset signal is high, set all the filter coefficient registers
34     // to zeros
35     // We're using a synchronous reset, which is recommended style for
36     // recent FPGA architectures
37     if(reset)begin
38         for(row = 0; row < FILTER_SIZE; row = row + 1) begin
39             for(col = 0; col < FILTER_SIZE; col = col + 1) begin
40                 r_f[row][col] <= 0;
41             end
42         end
43     // Otherwise, register the input filter coefficients into the 2D
44     // array signal
45     end else begin
46         for(row = 0; row < FILTER_SIZE; row = row + 1) begin
47             for(col = 0; col < FILTER_SIZE; col = col + 1) begin
48                 // Rearrange the 72-bit input into a 3x3
49                 // array of 8-bit filter coefficients.
50                 // signal[a +: b] is equivalent to
51                 // signal[a+b-1 : a]. You can try to plug
52                 // in
53                 // values for col and row from 0 to 2, to
54                 // understand how it operates.
55                 // For example at row=0 and col=0: r_f[0][0]
56                 // = i_f[0+:8] = i_f[7:0]
57                 // at row=0 and col=1:
58                 // r_f[0][1] = i_f[8+:8] = i_f[15:8]
59                 r_f[row][col] <= i_f[(row * FILTER_SIZE *
60                 // PIXEL_DATAW)+(col * PIXEL_DATAW) +:
61                 // PIXEL_DATAW];
```

```
51                     end
52             end
53         end
54     end*/
55
56
57
58 // Start of your code
59
60 localparam IMG_SIZE = 512 + 2; // size of the image plus padding(512 + 2)
61 localparam BUFFER_SIZE = IMG_SIZE*2 + FILTER_SIZE;
62 logic signed [PIXEL_DATAW-1:0] i_buffer [BUFFER_SIZE];
63
64 logic[19:0] pixel_counter; // counter to record how many pixels have been
  ↳ added to the buffer
65 logic[9:0] row_counter; // row counter to deal with the offset of adding
  ↳ padding into pixel_counter
66
67 logic valid_stage [0:4];
68
69 // output would be stored from multiplier and adder
70 logic signed [15:0] mult [0:8];
71 logic signed [18:0] add_1 [0:3];
72 logic signed [18:0] add_2 [0:1];
73 logic signed [18:0] add_3;
74
75 // The maximum bits from the last adder will be 20 bits
76 logic signed [19:0] add_final;
77
78 // capture output to 8 bits (0-255) grayscale
79 logic unsigned [7:0] cap_final;
80
81 // pipeline register to send and hold the value
82 logic signed [18:0] mult_reg [0:8];
83 logic signed [18:0] add_1_reg [0:3];
84 logic signed [18:0] add_2_reg [0:1];
85 logic signed [18:0] add_3_reg;
86 logic signed [18:0] add_wait [0:2]; // wait signal for add 1 to 3
87
88 always_ff @(posedge clk) begin
89     // If reset signal is high, set all the filter coefficient registers
  ↳ to zeros
```

```
90      // We're using a synchronous reset, which is recommended style for
91      // recent FPGA architectures
92      if(reset)begin
93
94          for(row = 0; row < FILTER_SIZE; row = row + 1) begin
95              for(col = 0; col < FILTER_SIZE; col = col + 1) begin
96                  r_f[row][col] <= 0;
97              end
98          end
99
100         for(buf_i = 0; buf_i < BUFFER_SIZE; buf_i = buf_i + 1)begin
101             i_buffer[buf_i] <= '0;
102         end
103
104         pixel_counter <= '0; // reset both counters
105         row_counter <= '0;
106         valid_stage[0] <= '0; // reset valid signal
107
108
109     // Otherwise, register the input filter coefficients into the 2D
110     // array signal
111     end else begin
112         for(row = 0; row < FILTER_SIZE; row = row + 1) begin
113             for(col = 0; col < FILTER_SIZE; col = col + 1) begin
114                 // Rearrange the 72-bit input into a 3x3
115                 // array of 8-bit filter coefficients.
116                 // signal[a +: b] is equivalent to
117                 // signal[a+b-1 : a]. You can try to plug
118                 // in
119                 // values for col and row from 0 to 2, to
120                 // understand how it operates.
121                 // For example at row=0 and col=0: r_f[0][0]
122                 // = i_f[0+:8] = i_f[7:0]
123                 // at row=0 and col=1:
124                 // r_f[0][1] = i_f[8+:8] = i_f[15:8]
125                 r_f[row][col] <= i_f[(row * FILTER_SIZE *
126                 // PIXEL_DATAW)+(col * PIXEL_DATAW) +:
127                 // PIXEL_DATAW];
128
129         end
130     end
131
132     // prepare for input buffer
```

```
122          // the newest input will be put input_buffer[0]
123          // rest will be shifted by 1
124          if(i_valid) begin
125              i_buffer[0] <= i_x;
126              pixel_counter <= pixel_counter +1;
127
128              for(buf_i = 0; buf_i < BUFFER_SIZE - 1; buf_i =
129                  ↵ buf_i + 1) begin
130                  i_buffer[buf_i+1] <= i_buffer[buf_i]; // 
131                  ↵ shift all pixels by 1 in the
132                  ↵ buffer
133
134          end
135
136          row_counter <= row_counter + 1; // to know which row
137          ↵ we are at
138
139          // sends the output from multipiler or adder to the
140          ↵ pipeline register
141
142          for (i = 0; i < 9; i++) begin // multipiler
143              mult_reg[i] <= mult[i];
144          end
145
146          for (i = 0; i < 4; i++) begin // adder 1
147              add_1_reg[i] <= add_1[i];
148          end
149          add_wait[0] <= mult_reg[0];
150
151          for (i = 0; i < 2; i++) begin // adder 2
152              add_2_reg[i] <= add_2[i];
153          end
154          add_wait[1] <= add_wait[0];
155
156          add_3_reg <= add_3; // adder 3
157          add_wait[2] <= add_wait[1];
158
159
160          // Capture the final output to 8 bits
161          if(add_final > 20'sd255) begin
162              cap_final <= 8'd255;
163          end else if(add_final < 20'sd0)begin
164              cap_final <= 8'd0;
```

```
159         end else begin
160             cap_final <= add_final[7:0];
161         end
162
163         // send the valid signal using pipeline register
164         valid_stage[1] <= valid_stage[0];
165         valid_stage[2] <= valid_stage[1];
166         valid_stage[3] <= valid_stage[2];
167         valid_stage[4] <= valid_stage[3];
168     end
169
170     // buffer for convolution
171     if(pixel_counter >= BUFFER_SIZE)begin
172         // addition output caused by two padding when
173         // → transfer
174         // rows stall valid signal for two cycles
175         if(row_counter == 'd1) begin
176             valid_stage[0] = 1'b0;
177         end else if(row_counter == 'd2)begin
178             valid_stage[0] = 1'b0;
179         end else begin
180             valid_stage[0] <= 1'b1;
181         end
182     end
183
184     if(row_counter == IMG_SIZE - 1)begin
185         // Reset row counter for new rows
186         row_counter <= 'd0;
187     end
188
189 end
190
191
192 mult8 mult_layer0_8 (.i_a(i_buffer[0]), .i_b(r_f[2][2]), .i_o(mult[8]));
193 mult8 mult_layer0_7 (.i_a(i_buffer[1]), .i_b(r_f[2][1]), .i_o(mult[7]));
194 mult8 mult_layer0_6 (.i_a(i_buffer[2]), .i_b(r_f[2][0]), .i_o(mult[6]));
195 mult8 mult_layer0_5 (.i_a(i_buffer[IMG_SIZE]), .i_b(r_f[1][2]),
196     → .i_o(mult[5]));
197 mult8 mult_layer0_4 (.i_a(i_buffer[IMG_SIZE + 1]), .i_b(r_f[1][1]),
198     → .i_o(mult[4]));
```

```
198 mult8 mult_layer0_3 (.i_a(i_buffer[IMG_SIZE + 2]), .i_b(r_f[1][0]),
  ↵ .i_o(mult[3]));
199 mult8 mult_layer0_2 (.i_a(i_buffer[IMG_SIZE*2]), .i_b(r_f[0][2]),
  ↵ .i_o(mult[2]));
200 mult8 mult_layer0_1 (.i_a(i_buffer[IMG_SIZE*2 + 1]), .i_b(r_f[0][1]),
  ↵ .i_o(mult[1]));
201 mult8 mult_layer0_0 (.i_a(i_buffer[IMG_SIZE*2 + 2]), .i_b(r_f[0][0]),
  ↵ .i_o(mult[0]));
202
203 add19 add_layer1_3 (.i_a(mult_reg[8]), .i_b(mult_reg[7]), .i_o(add_1[3]));
204 add19 add_layer1_2 (.i_a(mult_reg[6]), .i_b(mult_reg[5]), .i_o(add_1[2]));
205 add19 add_layer1_1 (.i_a(mult_reg[4]), .i_b(mult_reg[3]), .i_o(add_1[1]));
206 add19 add_layer1_0 (.i_a(mult_reg[2]), .i_b(mult_reg[1]), .i_o(add_1[0]));
207
208 add19 add_layer2_1 (.i_a(add_1_reg[3]), .i_b(add_1_reg[2]), .i_o(add_2[1]));
209 add19 add_layer2_0 (.i_a(add_1_reg[1]), .i_b(add_1_reg[0]), .i_o(add_2[0]));
210
211 add19 add_layer3_0 (.i_a(add_2_reg[1]), .i_b(add_2_reg[0]), .i_o(add_3));
212
213 add19 add_layer4_0 (.i_a(add_3_reg), .i_b(add_wait[2]), .i_o(add_final));
214
215 assign o_y = cap_final;
216 assign o_valid = valid_stage[4];
217 assign o_ready = i_ready;
218 // End of your code
219
220 endmodule
221
222 module mult8 (
223     input unsigned [7:0] i_a,
224     input signed [7:0] i_b,
225     output signed[15:0] i_o
226 );
227     assign i_o = $signed({1'b0, i_a}) * i_b;
228 endmodule
229
230 module add19(
231     input signed [18:0] i_a,
232     input signed [18:0] i_b,
233     output signed [19:0] i_o
234 );
235     assign i_o = i_a + i_b;
```

```
236 endmodule
237
238
```